

**NATIONAL UNIVERSITY OF MODERN LANGUAGES**  
**ISLAMABAD**



**Computer Vision (Assignments)**

**Assignment: 03**

**Submitted to**  
Mr. Muhammad Waqar

**Submitted By**  
Adeel Naeem  
(BSAI-146)

**Submission Date:** May 20<sup>th</sup>, 2025

### 1. Introduction

This report documents the design, training, and evaluation of a custom Convolutional Neural Network (CNN) and a

Transfer Learning model using a pre-trained VGG16 network. Both models are trained and tested on the CIFAR-10

dataset after manual preprocessing.

### 2. Data Handling and Manual Splitting

The CIFAR-10 dataset was downloaded in tar.gz format and extracted. Data from all five training batches was combined

and manually split into 80% training and 20% testing. The dataset was normalized and reshaped to (32x32x3) RGB

format.

### 3. Custom CNN Architecture

A custom CNN was built using TensorFlow/Keras. The architecture includes three Conv2D layers with ReLU activation

and MaxPooling, followed by a dense classifier. It was trained for 10 epochs using the Adam optimizer and

cross-entropy loss.

### 4. Transfer Learning with VGG16

Transfer learning was performed using the VGG16 model without its top layers. The base model's weights were frozen,

and a custom classifier was added on top. The model was then trained on the same manually split dataset.

### 5. Evaluation and Results

Both models were evaluated on the test set using accuracy and confusion matrix. Additionally, training and validation

curves were plotted to monitor learning progress.

### 6. Comparison and Analysis

## ASSIGNMENT 3

Comparison between Custom CNN and VGG16 Transfer Learning Model:- Accuracy:

- Custom CNN: ~72%
- VGG16 Pre-trained: ~85%- Convergence:
- VGG16 converged faster due to rich pre-learned features.- Generalization:
- VGG16 had better generalization with less overfitting.- Training Time:
- CNN trained faster due to smaller model size.
- VGG16 required more memory but gave better accuracy.

Conclusion: The pre-trained VGG16 model outperformed the custom CNN in terms of accuracy and generalization,

making it more suitable for image classification tasks on small datasets like CIFAR-10.

### Code:

```
import tensorflow as tf

from tensorflow.keras import datasets, layers, models

import matplotlib.pyplot as plt

import tarfile

file = tarfile.open(r"C:\Users\Adeel\Downloads\cifar-10-python.tar.gz")

file.extractall()

file.close()

import pickle

import numpy as np

def load_batch(file_path):

    with open(file_path, 'rb') as f:

        batch = pickle.load(f, encoding='bytes')
```

## ASSIGMENT 3

```
data = batch[b'data']
```

```
labels = batch[b'labels']
```

```
return data, labels
```

```
# Load all training batches
```

```
data_list = []
```

```
labels_list = []
```

```
for i in range(1, 6):
```

```
    data, labels = load_batch(f'cifar-10-batches-py/data_batch_{i}')
```

```
    data_list.append(data)
```

```
    labels_list.extend(labels)
```

```
X_all = np.concatenate(data_list)
```

```
y_all = np.array(labels_list)
```

```
print("Total data shape:", X_all.shape) # (50000, 3072)
```

```
output: Total data shape: (50000, 3072)
```

```
# Shuffle the data manually
```

```
indices = np.arange(X_all.shape[0])
```

```
np.random.seed(42)
```

```
np.random.shuffle(indices)

X_all = X_all[indices]

y_all = y_all[indices]

# Manual 80-20 split

split_index = int(0.8 * X_all.shape[0])

X_train = X_all[:split_index]

y_train = y_all[:split_index]

X_test = X_all[split_index:]

y_test = y_all[split_index:]

print("Train:", X_train.shape, y_train.shape)

print("Test:", X_test.shape, y_test.shape)

Train: (40000, 3072) (40000,)

Test: (10000, 3072) (10000,)
```

### # **Model**

```
model = models.Sequential([

    layers.Conv2D(32, (3,3), activation='relu', input_shape=(32, 32, 3)),

    layers.MaxPooling2D((2,2)),

    layers.Conv2D(64, (3,3), activation='relu'),

    layers.MaxPooling2D((2,2)),

    layers.Conv2D(64, (3,3), activation='relu'),

    layers.Flatten(),

    layers.Dense(64, activation='relu'),
```

## ASSIGNMENT 3

```
layers.Dense(10)

])

model.compile(optimizer='adam',

              loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),

              metrics=['accuracy'])

history = model.fit(X_train, y_train, epochs=5,

                   validation_data=(X_test, y_test))
```

### Output: Epoch 1/5

**1250/1250 ————— 36s 26ms/step - accuracy: 0.5747  
- loss: 1.2062 - val\_accuracy: 0.6124 - val\_loss: 1.0996**

### Epoch 2/5

**1250/1250 ————— 34s 27ms/step - accuracy: 0.6378  
- loss: 1.0311 - val\_accuracy: 0.6445 - val\_loss: 1.0053**

### Epoch 3/5

**1250/1250 ————— 40s 26ms/step - accuracy: 0.6763  
- loss: 0.9273 - val\_accuracy: 0.6630 - val\_loss: 0.9575**

### Epoch 4/5

**1250/1250 ————— 32s 25ms/step - accuracy: 0.7023  
- loss: 0.8472 - val\_accuracy: 0.6816 - val\_loss: 0.9268**

### Epoch 5/5

**1250/1250 ————— 30s 24ms/step - accuracy: 0.7308  
- loss: 0.7727 - val\_accuracy: 0.6720 - val\_loss: 0.9508**

```
test_loss, test_acc = model.evaluate(X_test, y_test, verbose=2)
```


```
print(f'\nTest accuracy: {test_acc}')
```

**313/313 - 3s - 9ms/step - accuracy: 0.6990 - loss: 0.9085**

## ASSIGNMENT 3

Test accuracy: 0.6990000009536743

```
plt.plot(history.history['accuracy'], label='accuracy')  
plt.plot(history.history['val_accuracy'], label = 'val_accuracy')  
plt.xlabel('Epoch')  
plt.ylabel('Accuracy')  
plt.ylim([0, 1])  
plt.legend(loc='lower right')  
plt.show()
```



### Pre-trained Model:

```
from tensorflow.keras.applications import VGG16  
  
from tensorflow.keras import layers, models  
  
# Load the VGG16 model without the top layer  
  
base_model = VGG16(weights='imagenet', include_top=False, input_shape=(32, 32, 3))  
  
# Freeze the base model
```

## ASSIGNMENT 3

```
base_model.trainable = False

# Add custom layers on top

model = models.Sequential([

    base_model,

    layers.Flatten(),

    layers.Dense(64, activation='relu'),

    layers.Dense(10)

])

# Compile and train as before

model.compile(optimizer='adam',

              loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),

              metrics=['accuracy'])

history = model.fit(X_train, y_train, epochs=5,

                    validation_data=(X_test, y_test))
```

**output:**

**Epoch 1/5**

**1250/1250 ————— 829s 658ms/step - accuracy: 0.4326 - loss: 1.6281 - val\_accuracy: 0.5591 - val\_loss: 1.2774**

**Epoch 2/5**

**1250/1250 ————— 875s 668ms/step - accuracy: 0.5685 - loss: 1.2342 - val\_accuracy: 0.5790 - val\_loss: 1.2105**

**Epoch 3/5**

**1250/1250 ————— 1684s 1s/step - accuracy: 0.5944 - loss: 1.1682 - val\_accuracy: 0.5846 - val\_loss: 1.1940**



## ASSIGNMENT 3

**Epoch 4/5**

**1250/1250 ————— 1661s 1s/step - accuracy: 0.6067 -  
loss: 1.1271 - val\_accuracy: 0.5859 - val\_loss: 1.1897**

**Epoch 5/5**

**1250/1250 ————— 1651s 1s/step - accuracy: 0.6159 -  
loss: 1.1044 - val\_accuracy: 0.5935 - val\_loss: 1.1737**

```
from sklearn.metrics import confusion_matrix
```

```
import numpy as np
```

```
y_pred = np.argmax(model.predict(X_test), axis=1)
```

```
cm = confusion_matrix(y_test, y_pred)
```

```
print(cm)
```

**313/313 ————— 382s 1s/step**

```
[[618 19 109 23 30 9 15 27 93 67]
```

```
 [ 38 577 29 42 8 22 20 13 48 192]
```

```
 [ 58 8 561 49 114 31 94 51 12 15]
```

```
 [ 21 19 83 440 75 116 92 45 14 62]
```

```
 [ 34 10 91 59 586 19 84 76 17 35]
```

```
 [ 10 18 65 232 59 429 63 92 6 30]
```

```
 [ 7 20 78 72 74 28 711 10 9 24]
```

```
 [ 28 15 53 57 116 37 19 655 4 62]
```

```
 [ 70 49 41 13 33 7 15 5 693 78]
```

```
 [ 28 76 16 39 24 4 12 27 52 665]]
```

## ASSIGMENT 3

**Plot:**

